

EXERCISE 2 RESULTS TABLE

1. This is how a schema is created.

The screenshot shows the Databricks workspace interface. On the left, the 'Catalog' pane displays a tree structure with 'fact_orders' and 'default' tables. The main pane shows the SQL editor with the following code:

```
1 CREATE SCHEMA IF NOT EXISTS my_coffee_shop
2 COMMENT 'Bright Coffee Shop analytics
3 schema';
4 USE my_coffee_shop;
5 SELECT current_database(); -- confirm:
6 should return my_coffee_shop
```

The results pane below the code shows a table with one row:

current_database()
my_coffee_shop

2. Benefits of creating this table is to track business product, for analysis, which products are sold fast, products, prices etc.

The screenshot shows the Databricks workspace interface. On the left, the 'Catalog' pane displays a tree structure with 'fact_orders' and 'default' tables. The main pane shows the SQL editor with the following code:

```
1 from my_coffee_shop.dim_product_table
```

The results pane below the code shows a table with 10 rows:

product_id	product_code	product_name	price	category
2	CAP001	Cappuccino	32.50	Coffee
3	TEA001	Green Tea	18.00	Tea
4	TEA002	Black Tea	15.50	Tea
5	CAKE001	Cheesecake Slice	45.00	Food
6	ESP001	Espresso	25.00	Coffee
7	CAP001	Cappuccino	32.50	Coffee
8	TEA001	Green Tea	18.00	Tea
9	TEA002	Black Tea	15.50	Tea
10	CAKE001	Cheesecake Slice	45.00	Food

3. Creating dim_product table columns, populating the columns with correct data, and using the select * to run the whole table results

dbc-7e8178e3-3b70.cloud.databricks.com/editor/notebooks/4317377202980985?o=7474652257519875#command/6875888060417394

workspace

Exercise two Data Definition Language & Dat...

1 alter table my_coffee_shop.
dim_table_customers
2 rename column loyalty points to points

See performance (1) Optimize

06:55 PM (23s) 13

1 select *
2 from my_coffee_shop.dim_table_customers

See performance (1) Optimize

customer_id	first_name	last_name	email	tier	poir
-------------	------------	-----------	-------	------	------

Exercise two Data Definition Language & Dat...

07:29 PM (18s) 15

Go to last run cell

5 (2, 'tsholo', 'mokoena', 't
com', 'gold'),
6 (3, 'kagiso', 'banda', 'kag
com', 'silver'),
7 (4, 'kabelo', 'moema', 'kab
com', 'silver'),
8 (5, 'siphon', 'moroke', 'sip
com', 'bronze');

See performance (1)

num_affected_rows	num_inserted_rows
1	5

Cell 1
Cell 2
Cell 3
Cell 5
Cell 6
Cell 8
Cell 10
Cell 12
Cell 13
Cell 15
Cell 16
Cell 18
Cell 19
Cell 20

Exercise two Data Definition Language & Dat...

07:30 PM (4s) 16

1 select *
2 from my_coffee_shop.dim_table_customers

See performance (1) Optimize

customer_id	first_name	last_name	email	tier
1	mpho	tsolo	mpho@gmail.com	gold
2	tsholo	mokoena	tsholo@gmail.com	gold
3	kagiso	banda	kagiso@gmail.com	silver
4	kabelo	moema	kabelo@gmail.com	silver

5 rows | 3.51s runtime Refreshed 4 hours ago

This result is stored as _sqlidf and can be used in other Python and SQL cells. Send feedback

4. Using tblproperties for column mapping, inserting values into the column and running the code for results.

The first screenshot shows the Databricks workspace interface. The left sidebar contains navigation options like Home, Learn, Workspace, Recents, Catalog, Jobs & Pipelines, Compute, Discover, Marketplace, SQL, SQL Editor, Queries, Dashboards, Genie Spaces, Alerts, Query History, and SQL Warehouses. The main panel displays a SQL query editor with the following code:

```
1 alter table my_coffee_shop.  
  dim_product_table  
2 set TBLProperties ('delta.columnmapping.  
  mode'='name')
```

The query results show 5 rows and 3.51s runtime. A message indicates: "This result is stored as _sqlidf and can be used in other Python and SQL cells." Below the query, a question is posed: "Question 6 insert atleast 5 products into dim_products".

The second screenshot shows the same workspace after executing a query. The SQL editor displays the following code:

```
1 select *  
2 from my_coffee_shop.dim_product_table
```

The query results show 10 rows and 2.85s runtime. The results are displayed in a table format:

product_id	product_code	product_name	price	category
1	ESP001	Espresso	25.00	Coffee
2	CAP001	Cappuccino	32.50	Coffee
3	TEA001	Green Tea	18.00	Tea
4	TEA002	Black Tea	15.50	Tea
5	CAKE001	Cheesecake Slice	45.00	Food
6	ESP001	Espresso	25.00	Coffee
7	CAP001	Cappuccino	32.50	Coffee
8	TEA001	Green Tea	18.00	Tea
9	TEA002	Black Tea	15.50	Tea
10	CAKE001	Cheesecake Slice	45.00	Food

5. Backup database results

The screenshot shows the Databricks workspace interface. On the left, the 'Catalog' sidebar is open, displaying a tree view of databases and tables. The 'fact_orders' table is selected, showing its schema: order_id, customer_id, product_id, store_id, date_id, quantity, unit_price, and total_amount. The main panel displays a SQL query titled 'Exercise two Data Definition Language & Dat...'. The query is as follows:

```
1 ----verifying the table
2 select *
3 from my_coffee_shop.
   dim_table_customers_version1_backup
```

The query results are displayed in a table with 5 rows and 6 columns: customer_id, first_name, last_name, email, tier, and an unnamed column. The data is as follows:

	customer_id	first_name	last_name	email	tier
1	1	mpho	tsolo	mpho@gmail.com	gold
2	2	tsholo	mokoena	tsholo@gmail.com	gold
3	3	kagiso	banda	kagiso@gmail.com	silver
4	4	kabelo	moema	kabelo@gmail.co...	silver
5					

The query was executed at 10:49 PM (2s) and returned 5 rows with a runtime of 2.13s. The results were refreshed 39 minutes ago.

6. Describe history table

The screenshot shows the Databricks workspace interface. On the left, the 'Catalog' sidebar is open, displaying a tree view of databases and tables. The 'fact_orders' table is selected, showing its schema: order_id, customer_id, product_id, store_id, date_id, quantity, unit_price, and total_amount. The main panel displays a SQL query titled 'Exercise two Data Definition Language & Dat...'. The query is as follows:

```
1 --- confirming the deleted column
2
3 DESCRIBE HISTORY my_coffee_shop.
   dim_table_customers_version1_backup
```

The query results are displayed in a table with 3 rows and 6 columns: version, timestamp, user_id, user_name, operation, and an unnamed column. The data is as follows:

	version	timestamp	user_id	user_name	operation
1	2	2026-05-28T20:54:11.000+00...	735425528238...	mokoneprecious@gmail.co...	OPTIMIZ
2	1	2026-05-28T20:54:08.000+00...	735425528238...	mokoneprecious@gmail.co...	DELETE
3	0	2026-05-28T20:45:46.000+00...	735425528238...	mokoneprecious@gmail.co...	CREATE

The query was executed at 10:56 PM (4s) and returned 3 rows with a runtime of 3.61s. The results were refreshed 33 minutes ago.